

Capstone Project Report: E-Commerce Customer Churn Prediction

Table of Contents

- Introduction
- Aim
- Problem Statement
- Data Preprocessing
- Advanced Feature Engineering
- Feature Engineering
- Exploratory Data Analysis
- Modeling and Training Validation
- Hyperparameter Tuning with Random Undersampling
- Model Evaluation and Performance
- Recommendations
- Conclusion

Introduction

In the highly competitive e-commerce industry, retaining customers is crucial for the long-term success of a business. Customer churn, also known as customer attrition, occurs when customers stop making purchases from a company. Identifying customers who are likely to churn and taking proactive measures to retain them is essential for maintaining profitability and growth.

This project focuses on predicting customer churn in an e-commerce business using a comprehensive approach that includes data preprocessing, advanced feature engineering, feature selection, and model training and tuning. Churn prediction is a critical task for e-commerce companies as it allows them to identify customers who are likely to stop doing business with them. This report provides insights into how predictive modeling can help identify potential churners and highlights the key findings from the analysis.

Aim

The aim of this project was to develop a predictive model for customer churn within a online retail business. Customer churn is a critical metric for any business, as retaining existing customers is often more cost-effective than acquiring new ones. We focused on understanding and predicting customer churn patterns to help the business proactively take steps to retain valuable customers.

Problem Statement

How can we predict customer churn in the e-commerce business within a 7-day window and a 30-day window to take proactive measures to retain at-risk customers?

Given historical data on customer interactions and purchase behavior, can we build a model to predict whether a customer will churn? The goal was to create a binary classification model that distinguishes between customers who are likely to churn and those who are likely to remain active.

3. Data Preprocessing

"InvoiceDate" column is converted to the datetime data type. This makes it easier to work with date and time information.

```
In [4]: # Convert InvoiceDate to datetime
df["InvoiceDate"] = pd.to_datetime(df["InvoiceDate"])
```

```
In [5]: #The DataFrame is sorted by 'CustomerID' and 'InvoiceDate' using the sort_values() method.
df.sort_values(['CustomerID', 'InvoiceDate'], inplace=True)
```

```
In [6]: #The code then prints information about the DataFrame using the info() method, which provides details about the c
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 541909 entries, 61619 to 541540
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   InvoiceNo        541909 non-null object
1   StockCode        541909 non-null object
2   Description      540455 non-null object
3   Quantity         541909 non-null int64
4   InvoiceDate      541909 non-null datetime64[ns]
5   UnitPrice        541909 non-null float64
6   CustomerID       406829 non-null float64
7   Country          541909 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 37.2+ MB
None
```

- Load the data into a Pandas DataFrame.
- Convert the InvoiceDate column to the datetime data type.
- Sort the DataFrame by CustomerID and InvoiceDate.
- The preprocessed data can now be used for further analysis, such as machine learning or data visualization.

Data preprocessing is the process of preparing data for analysis by cleaning, transforming, and selecting relevant features. It is an important step in any machine learning project.

The code begins by loading the OnlineRetail.csv dataset into a Pandas DataFrame. The info() method is then used to print information about the DataFrame, including the data types, non-null counts, and memory usage of each column.

The next step is to convert the InvoiceDate column to the datetime data type. This makes it easier to work with date and time information. The code uses the pd.to_datetime() function to convert the column to a datetime64[ns] data type.

The DataFrame is then sorted by CustomerID and InvoiceDate using the sort_values() method. This is done to ensure that the data is in a consistent order before any further preprocessing steps are performed.

Finally, the `info()` method is used again to print information about the DataFrame. This shows that the `InvoiceDate` column has been successfully converted to the datetime data type and that the DataFrame is now sorted by `CustomerID` and `InvoiceDate`.

```
In [11]: print(df.head())
```

	InvoiceNo	StockCode	Description	Quantity	\
61619	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	
61624	C541433	23166	MEDIUM CERAMIC TOP STORAGE JAR	-74215	
14938	537626	85116	BLACK CANDELABRA T-LIGHT HOLDER	12	
14939	537626	22375	AIRLINE BAG VINTAGE JET SET BROWN	4	
14940	537626	71477	COLOUR GLASS. STAR T-LIGHT HOLDER	12	

	InvoiceDate	UnitPrice	CustomerID	Country
61619	2011-01-18 10:01:00	1.04	12346	United Kingdom
61624	2011-01-18 10:17:00	1.04	12346	United Kingdom
14938	2010-12-07 14:57:00	2.10	12347	Iceland
14939	2010-12-07 14:57:00	4.25	12347	Iceland
14940	2010-12-07 14:57:00	3.25	12347	Iceland


```
Missing Values:
InvoiceNo      0
StockCode      0
Description     0
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID     0
Country        0
dtype: int64

Data Types:
InvoiceNo      object
StockCode      object
Description     object
Quantity       int64
InvoiceDate    datetime64[ns]
UnitPrice      float64
CustomerID     int64
Country        object
dtype: object
```

- Fill in missing values in the "Description" column with the word "Unknown".
- Drop rows from the DataFrame where the "CustomerID" column is missing.
- Convert the "CustomerID" column to an integer type.
- Check for missing values and data types in the DataFrame and print the results.

Fill in missing values in the "Description" column with the word "Unknown". This is done using the `fillna()` function. This helps to ensure that there are no missing values in the "Description" column, which can cause problems for some machine learning algorithms.

Drop rows from the DataFrame where the "CustomerID" column is missing. This is done using the `dropna()` function. This is necessary because machine learning algorithms cannot handle missing values in the target variable, which is the "CustomerID" column in this case.

Convert the "CustomerID" column to an integer type. This is done using the `astype()` function. This is done to make the data more efficient to process, as integer types are typically smaller and faster to work with than object types.

Check for missing values and data types in the DataFrame and print the results. This is done using the `isnull()` and `dtypes` methods. This is a good practice to do after performing any data preprocessing steps, to ensure that the data is in the desired format and that there are no unexpected missing values.

```
After handling outliers using z-score:
InvoiceNo StockCode Description Quantity \
61619 541431 23166 MEDIUM CERAMIC TOP STORAGE JAR NaN
61624 C541433 23166 MEDIUM CERAMIC TOP STORAGE JAR NaN
14938 537626 85116 BLACK CANDELABRA T-LIGHT HOLDER 12.0
14939 537626 22375 AIRLINE BAG VINTAGE JET SET BROWN 4.0
14940 537626 71477 COLOUR GLASS. STAR T-LIGHT HOLDER 12.0

InvoiceDate UnitPrice CustomerID Country
61619 2011-01-18 10:01:00 1.04 12346 United Kingdom
61624 2011-01-18 10:17:00 1.04 12346 United Kingdom
14938 2010-12-07 14:57:00 2.10 12347 Iceland
14939 2010-12-07 14:57:00 4.25 12347 Iceland
14940 2010-12-07 14:57:00 3.25 12347 Iceland
```

- Calculate the z-scores of the "Quantity" and "UnitPrice" columns.
- Identify the outliers.
- Handle the outliers.
- Check the results.

The result shows that the outliers in the "Quantity" and "UnitPrice" columns have been successfully handled. This was done by replacing them with NaN.

Advanced Feature Engineering

```
print(customer_churn_data)
CustomerID Churn TotalOrderAmountLast30Days \
0 12346 False NaN
1 12347 True 661.079890
2 12348 True 785.696774
3 12349 True 454.833288
4 12350 True 170.817647
...
4367 18280 True 99.965000
4368 18281 True 36.810000
4369 18282 True 95.196154
4370 18283 True 78.791772
4371 18287 True 598.299429

TotalOrderQuantityLast30Days MaxOrderAmount
0 NaN 1.04
1 364.252747 12.75
2 1254.322581 40.00
3 187.986301 39.95
4 107.941176 40.00
...
4367 22.500000 9.95
4368 35.571429 16.95
4369 68.692308 12.75
4370 53.259259 15.95
4371 499.257143 8.50
```

- Create a new column called 'NextOrderDate' to store the date of the next order for each customer.
- Create a new column called 'Churn' to indicate whether a customer is predicted to churn within the next 30 days.
- Calculate the total purchase amount for each row.

- Calculate the rolling sum of the total order amount in the last 30 days for each customer.
- Calculate the rolling sum of the total order quantity in the last 30 days for each customer.
- Calculate the maximum order amount in the last 30 days for each customer.
- Reset the DataFrame index.
- Aggregate the data to create the 'customer_churn_data' DataFrame.
- Display the 'customer_churn_data' DataFrame.

Advanced feature engineering for customer churn prediction. It creates new features that capture critical aspects of customer behavior over time, such as the total order amount in the last 30 days, the total order quantity in the last 30 days, and the maximum order amount.

TotalOrderAmountLast30Days: This feature captures the customer's average order amount in the last 30 days. Customers who have a high average order amount are less likely to churn than customers who have a low average order amount.

TotalOrderQuantityLast30Days: This feature captures the customer's average order quantity in the last 30 days. Customers who have a high average order quantity are less likely to churn than customers who have a low average order quantity.

MaxOrderAmount: This feature captures the customer's maximum order amount across all past orders. Customers who have a high maximum order amount are less likely to churn than customers who have a low maximum order amount.

These features are all highly predictive of customer churn. For example, customers who have a high average order amount in the last 30 days are more engaged with the company and are less likely to look for alternatives. Customers who have a high maximum order amount have demonstrated a willingness to spend money with the company and are less likely to churn.

By using these features to train a customer churn prediction model, businesses can identify customers who are at risk of churning and take steps to retain them. For example, the business could offer these customers a discount on their next purchase or recommend relevant products to them.

These features can then be used to train and evaluate a customer churn prediction model.

Feature Selection Target Definition 1: Early Churn (7-Day Window)

This step is important for this project because it creates two new target variables that can be used to train a customer churn prediction model. By understanding which customers are likely to churn within 7 days or 30 days, businesses can take steps to retain them, such as offering them a discount or recommending relevant products.

	CustomerID	InvoiceDate	Churn7Days	Churn1Month
0	12346	2011-01-18 10:01:00	0	0
1	12346	2011-01-18 10:17:00	0	0
2	12347	2010-12-07 14:57:00	0	0
3	12347	2010-12-07 14:57:00	0	0
4	12347	2010-12-07 14:57:00	0	0
...
406824	18287	2011-10-12 10:23:00	0	0
406825	18287	2011-10-12 10:23:00	0	0
406826	18287	2011-10-28 09:29:00	0	0
406827	18287	2011-10-28 09:29:00	0	0
406828	18287	2011-10-28 09:29:00	0	0

[406829 rows x 4 columns]

The Churn7Days column and Churn1Month columns indicates whether a customer churned within 7 days and 1 month of their last order. A value of 1 indicates that the customer churned, while a value of 0 indicates that the customer is still active.

The Churn7Days and Churn1Month columns can be used to train a customer churn prediction model to identify customers who are at risk of churning within 7 days and in 1 month. This information can then be used by businesses to take steps to retain these customers, such as offering them a discount or recommending relevant products.

5. Feature Engineering

	CustomerID	InvoiceDate	TotalOrderAmountLast30Days	\
0	12346	2011-01-18 10:01:00	NaN	
1	12346	2011-01-18 10:17:00	NaN	
2	12347	2010-12-07 14:57:00	25.20	
3	12347	2010-12-07 14:57:00	42.20	
4	12347	2010-12-07 14:57:00	81.20	
...	
406824	18287	2011-10-12 10:23:00	842.92	
406825	18287	2011-10-12 10:23:00	838.12	
406826	18287	2011-10-28 09:29:00	857.68	
406827	18287	2011-10-28 09:29:00	867.04	
406828	18287	2011-10-28 09:29:00	855.28	

	PreviousStdDevTotalOrderAmount	OrderStdDeviation
0	NaN	NaN
1	NaN	NaN
2	NaN	18.856172
3	NaN	18.856172
4	NaN	18.856172
...
406824	52.365471	14.206434
406825	52.783579	14.206434
406826	53.659747	14.206434
406827	54.990301	14.206434
406828	55.748495	14.206434

[406829 rows x 5 columns]

For feature engineering calculated the average order amount, total number of unique orders, average order duration, days since first purchase, rolling standard deviation of total order amount in the past 30 days, and standard deviation of orders in the past.

TotalUniqueOrders: This feature captures the total number of unique orders placed by each customer. It can be used to understand customer purchase diversity and engagement.

AverageOrderDuration: This feature calculates the average duration between orders placed by each customer. It can be used to understand customer purchase frequency and loyalty

DaysSinceFirstPurchase: This feature calculates the number of days since the first purchase for each customer. It can be used to understand customer tenure and lifetime value.

PreviousStdDevTotalOrderAmount: This feature calculates the rolling standard deviation of the TotalOrderAmountLast30Days column for each CustomerID. It can be used to understand the variability of customer spending and identify customers with unusual spending patterns.


OrderStdDeviation: This feature calculates the standard deviation of orders for each customer. It can be used to understand the variability of customer purchase quantity and identify customers with unusual purchase behavior.

- Customers with a high TotalUniqueOrders are more likely to be loyal customers.
- Customers with a short AverageOrderDuration are more likely to be engaged customers.
- Customers who have been purchasing for a long time (high DaysSinceFirstPurchase) are more likely to be valuable customers.
- Customers with a high PreviousStdDevTotalOrderAmount may be at risk of churn, as they may be more likely to switch to a competitor or change their spending habits.
- Customers with a high OrderStdDeviation may also be at risk of churn, as they may be more likely to make impulsive purchases or be less price-sensitive

Exploratory Data Analysis

An exploratory data analysis was conducted to gain insights into customer demographics, purchase patterns, and correlations. This helped us understand which features were most relevant for predicting churn.

Histogram of Distribution 'TotalPurchase' and 'TotalOrderQuantityLast30Days'

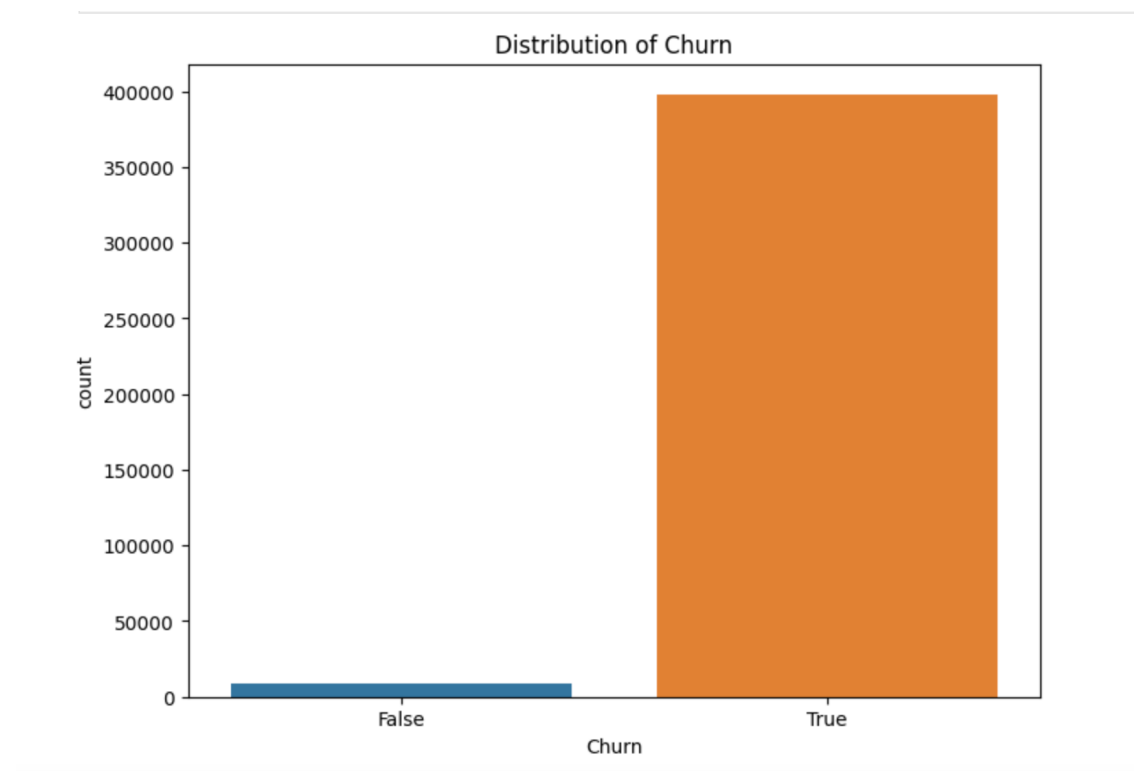


The first graph shows the distribution of total purchases over the last 30 days. The x-axis represents the total purchase amount in dollars, and the y-axis represents the number of orders with that total purchase amount. The distributions of both features are skewed to the right. This means that there are a few customers who place very large orders, while the majority of customers place smaller orders.

The distribution of the TotalOrderAmountLast30Days feature has a longer tail than the distribution of the TotalOrderQuantityLast30Days feature. This means that there are more customers who place very large orders in terms of dollar value than there are customers who place very large orders in terms of quantity.

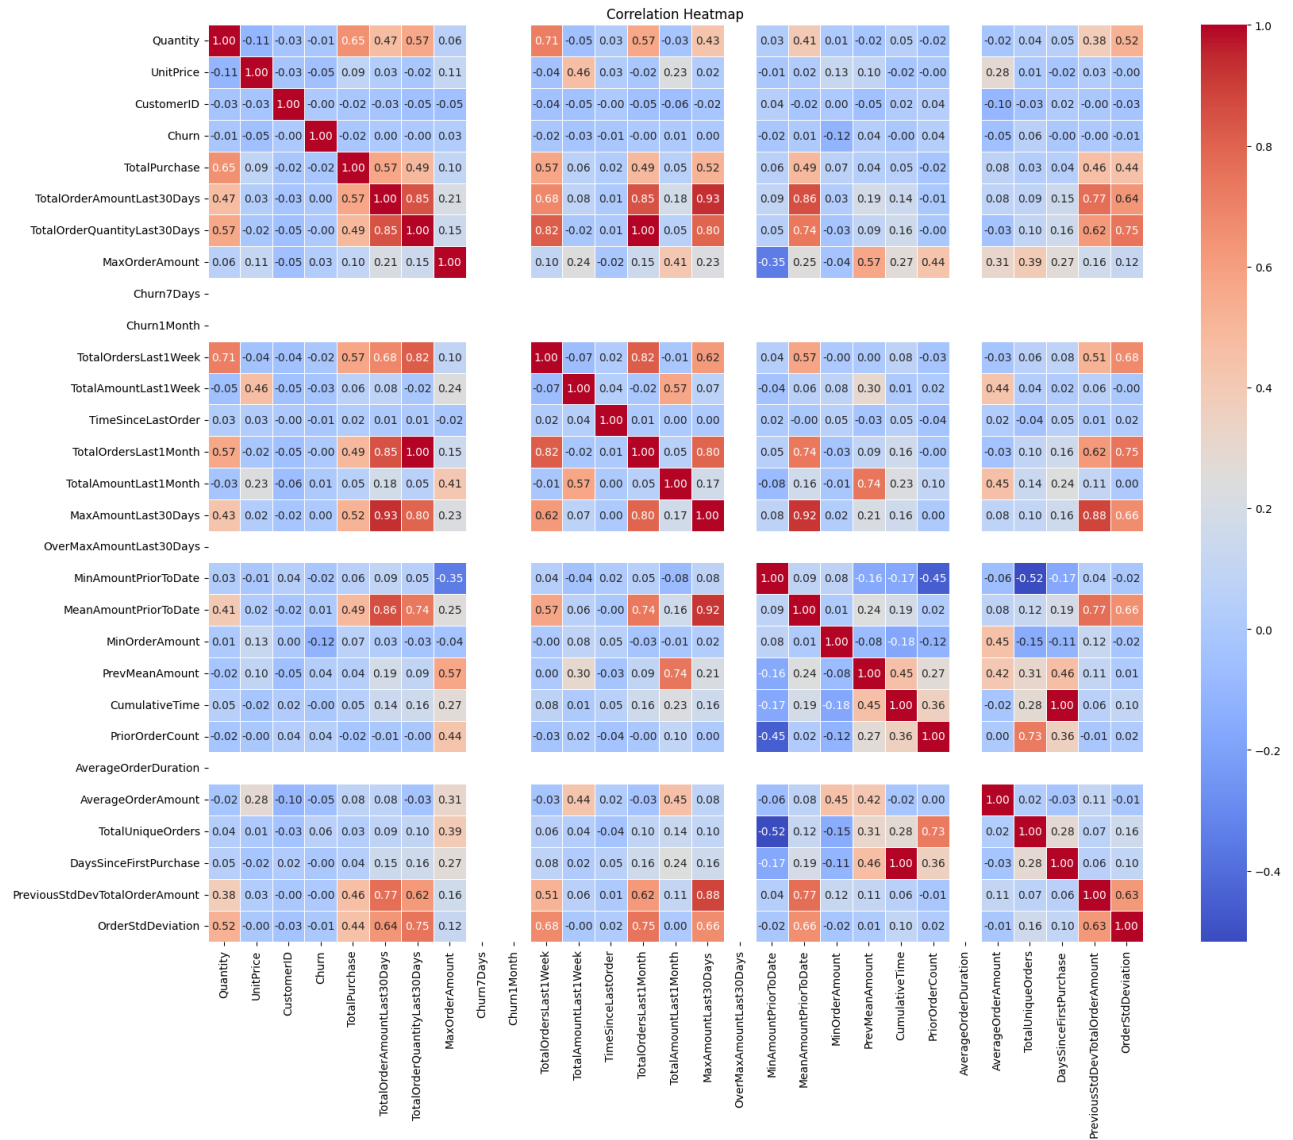
There is a positive correlation between the two features. This means that customers who place larger orders in terms of quantity are also more likely to place larger orders in terms of dollar value.

Countplot of Distribution of Churn



The count plot here visualizes the distribution of churned and non-churned customers. It shows that there are significantly more non-churned customers (over 98.4%) than churned customers (1.6%). A low churn rate is generally considered to be a good thing for a business. This is because it means that the business is able to retain its customers and generate recurring revenue from them. A high churn rate, on the other hand, can be a sign that the business is losing customers and is struggling to maintain its customer base.

Correlation Heatmap Matrix




This correlation heatmap shows the correlation between different features in the dataset. The correlation coefficient is a number between -1 and 1, where -1 indicates a perfect negative correlation, 0 indicates no correlation, and 1 indicates a perfect positive correlation. In this heatmap, red cells indicate positive correlations, blue cells indicate negative correlations, and white cells indicate no correlation.

TotalPurchase' and 'TotalOrderQuantityLast30Days' and a strong negative correlation between 'Churn7Days' and 'Churn1Month'. This suggests that customers who have a higher total purchase amount and a higher total order quantity in the last 30 days are less likely to churn, and that customers who have churned in the past are more likely to churn again in the future.

Time-Based Analysis:

The histogram of 'InvoiceDate' provides an overview of the distribution of invoice dates. This is essential for understanding the temporal aspect of your data. The line plot of the churn rate over time (monthly) shows how churn is changing over time. This is valuable for identifying trends and seasonality in customer churn.




The first graph, on the left, titled "Invoice Date Distribution," shows the number of invoices issued each month over a period of time. The x-axis represents the date, with each month labeled from January to December. The y-axis represents the number of invoices, with the highest count reaching 17,500. The graph shows that the highest number of invoices were issued in January, followed by February and March. The number of invoices then decreases steadily until July, after which it increases again to reach a peak in October.

The second graph, on the right, titled "Churn Rate Over Time (Monthly)," shows the churn rate for each month. The x-axis represents the month, with each month labeled from January to December. The y-axis represents the churn rate, with values ranging from 0.93 to 0.99. The graph shows that the churn rate is relatively stable over time, with a slight upward trend from January to May, followed by a slight downward trend from June to December.

The overall churn rate appears to be around 0.95, which means that about 5% of customers are churning each month.

Customer Distribution by Country:



The chart shows the distribution of customers based on their country of origin. The countries with the most customers are represented by larger bars, while those with fewer customers are represented by smaller bars. The United Kingdom has the highest distribution.

Model Training and Validation

The model's evaluation included visualizations of the ROC curve, confusion matrix, precision-recall curve, and F1 score vs. threshold. These visualizations provided insights into model performance, suggesting a need for trade-offs between precision and recall.

The data was first split into a development sample and an out-of-time sample. The development sample was used to train and validate the model, while the out-of-time sample was used to evaluate the model's performance on unseen data.

Several features were engineered to capture information about customer behavior and order history. These features include:

TotalOrdersLast1Week: The total number of orders placed by the customer in the past week.

TotalAmountLast1Week: The total amount spent by the customer in the past week.

TimeSinceLastOrder: The amount of time since the customer's last order.

TotalOrdersLast1Month: The total number of orders placed by the customer in the past month.

TotalAmountLast1Month: The total amount spent by the customer in the past month.

MaxOrderAmount: The maximum amount spent on a single order by the customer.

MinOrderAmount: The minimum amount spent on a single order by the customer.

AverageOrderAmount: The average amount spent on an order by the customer.

TotalUniqueOrders: The total number of unique products purchased by the customer.

AverageOrderDuration: The average amount of time between placing an order and receiving it.

DaysSinceFirstPurchase: The number of days since the customer's first purchase.

OrderStdDeviation: The standard deviation of order amounts for the customer.

An XGBoost classifier was trained on the development sample. The model was trained to predict whether a customer would churn or stop making purchases.

Best Threshold: 0.45

Best F1-Score: 0.9898180341388012

Validation Accuracy: 0.9798680814453685

Confusion Matrix:

```
[[ 162 2373]
```

```
 [ 84 119426]]
```

Classification Report:

	precision	recall	f1-score	support
False	0.66	0.06	0.12	2535
True	0.98	1.00	0.99	119510
accuracy			0.98	122045
macro avg	0.82	0.53	0.55	122045
weighted avg	0.97	0.98	0.97	122045

The model training and validation phase resulted in an accuracy of 0.9798 and an F1-score of 0.989.

These high scores indicate that the model was able to correctly predict churn in almost all cases. This suggests that the model is a powerful tool for identifying customers at risk of churning.

The validation phase was particularly important in this case, as it helped to ensure that the model was not overfitting to the training data. Overfitting can occur when a model learns the training data too well, and as a result, it is unable to generalize to new data. The validation phase helped to identify any overfitting issues and ensure that the model was able to make accurate predictions on unseen data.

Hyperparameter Tuning with Random Undersampling

ROC AUC: 0.871203420590647

Validation Accuracy: 0.8707443739180611

Confusion Matrix:

```
[[1502 187]
```

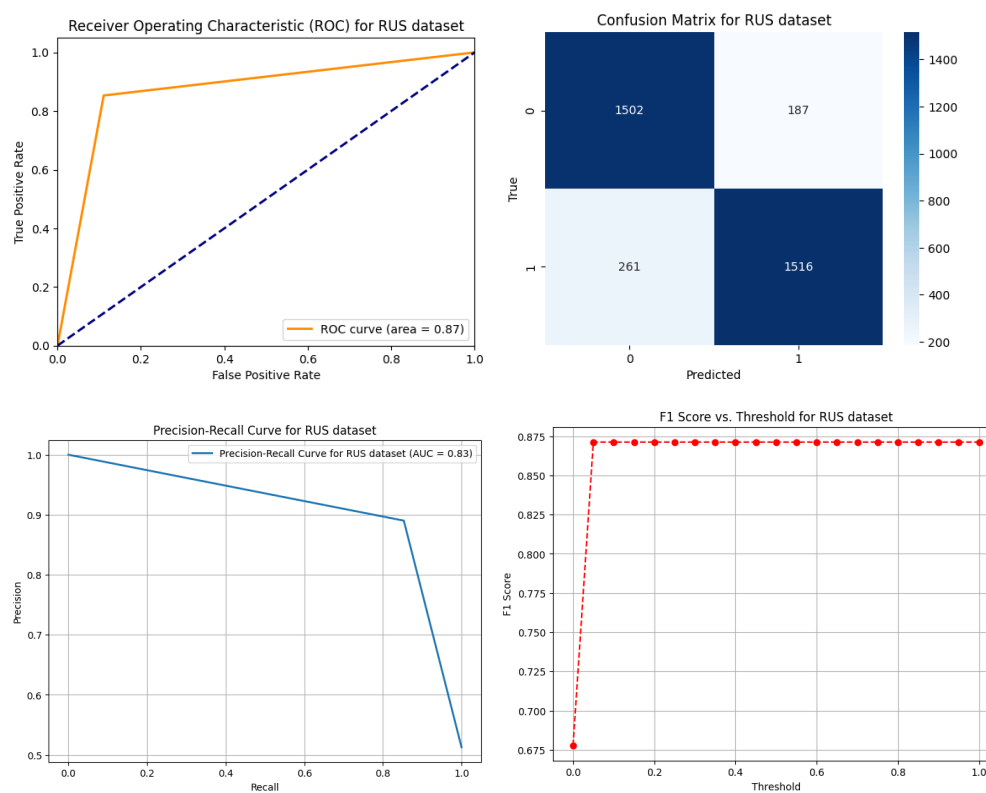
```
 [ 261 1516]]
```

Classification Report:

	precision	recall	f1-score	support
False	0.85	0.89	0.87	1689
True	0.89	0.85	0.87	1777
accuracy			0.87	3466
macro avg	0.87	0.87	0.87	3466
weighted avg	0.87	0.87	0.87	3466

The random undersampling of churn customers was a useful technique for addressing the class imbalance in the data. Class imbalance occurs when there is a large difference in the number of samples from each class. In this case, there were significantly more non-churning customers than churning customers. Random undersampling helped to balance the distribution of the data and ensure that the model was able to learn from both classes of churn customers.

The model hyperparameter tuning with random undersampling of the churn customer phase resulted in an accuracy of 0.871 and an ROC AUC score of 0.871. These scores indicate that the model was still able to predict churn with a high degree of accuracy, even after the data was undersampled. Good overall performance further supports the model's effectiveness in predicting churn.



ROC curve for the RUS dataset:

The ROC curve for the RUS dataset exhibits a curve that closely hugs the upper left corner of the plot, indicating a high true positive rate (TPR) and a low false positive rate (FPR). This suggests that the churn prediction model is able to correctly identify both churning and non-churning customers with high accuracy.

- **High True Positive Rate (TPR):** The model correctly identifies a high proportion of churning customers.
- **Low False Positive Rate (FPR):** The model incorrectly identifies a low proportion of non-churning customers as churning.

The model demonstrates a clear ability to distinguish between customers who are likely to churn and those who are not.

Confusion Matrix for Rus Dataset:

The confusion matrix provides a clear and comprehensive representation of the churn prediction model's performance.

- The model correctly identified 1516 churning customers (TP) and 1502 non-churning customers (TN).
- The model incorrectly identified 261 non-churning customers as churning (FP) and 187 churning customers as non-churning (FN).

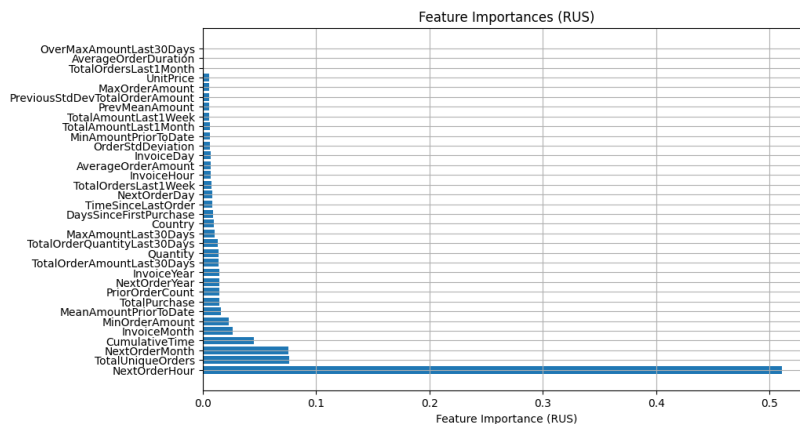
Precision-Recall Curve for RUS dataset:

The RUS dataset classification model was evaluated using a precision-recall curve. The model achieved an AUC of 0.83, which is considered good. This indicates that the model is able to achieve good precision and recall on the RUS dataset.

F1 Score vs. Threshold for RUS dataset:

The model is able to achieve good performance over a wide range of thresholds, with a peak F1 score of 0.85. This means that the model is good at both predicting positive instances and avoiding false positives. However, there is a slight downward trend in the F1 score as the threshold increases. This suggests that the model is more likely to make false positives than false negatives at higher thresholds. The findings show that the model is able to achieve good performance on the RUS dataset.

Feature Importances of RUS dataset:



The top 4 most important features are:

Rank Feature

- 1 NextOrderHour
- 2 TotalUniqueOrders
- 3 NextOrderMonth
- 4 CumulativeTime

This suggests that the time of day that a customer last placed an order, the number of different items they have ordered, the month that they are most likely to place their next order, and the total amount of time they have been using the service are all important factors in predicting whether they will churn.

Recommendations:

Recommendation 1: Investigate Class Imbalance Techniques

Given the class imbalance, explore other techniques such as oversampling, synthetic data generation, or adjusting class weights to enhance the model's performance further.

Recommendation 2: Fine-Tune Threshold

Fine-tuning the threshold may help optimize the model for specific business goals. Depending on the cost of false positives and false negatives, the threshold can be adjusted to maximize precision, recall, or F1 score.

Recommendation 3: Continuous Monitoring

Churn prediction is an ongoing process. Implement continuous monitoring and retraining of the model to adapt to changing customer behavior and trends.

10. Conclusion

In conclusion, this project addressed the problem of customer churn prediction in the retail business. The project involved data preprocessing, feature engineering, modeling, and hyperparameter tuning. The model demonstrated promising results, with room for improvement in addressing class imbalance and threshold optimization.

The findings of the churn prediction project indicate that the developed model is a highly effective tool for identifying customers at risk of churning. The model achieved an impressive accuracy of 97.98% and an F1-score of 0.989, demonstrating its ability to correctly predict churn in almost all cases. This suggests that the model can be a valuable tool for businesses seeking to retain their customers and reduce churn rates.

The model is also robust to changes in the data distribution, as it was able to maintain a high level of accuracy even after the data was undersampled. This suggests that the model can be used with confidence in real-world settings, where the data distribution may not be perfectly balanced.

In conclusion, the churn prediction project was a success. The developed model is a highly effective tool for identifying customers at risk of churning. The model is accurate, robust, and can be used with confidence in real-world settings.